Introduction

Richèl J.C. Bilderbeek[1]

[1]Groningen Institute for Evolutionary Life Sciences, University of Groningen, Groningen, The Netherlands

September 27, 2019

1

INTRODUCTION

S PECIATION is the process that creates new species, connecting all of life to one shared
common ancestor. It is a process that has resulted in the millions of species on Earth
nowadays, as well as in the many species that have gone extinct. Some speciation events
happened earlier than others, from hundreds of millions of years ago (so-called 'long-
enduring species', or, informally, 'living fossil') to more recent ones. See figure 1.1 shows
an example of each.

Figure 1.1 | An long-enduring species (left) and a young species (right). The species at the left is a preserved
specimen of *Latimeria chalumnae*, estimated to exist for hundreds of millions of year. The species at the right is
the *Homo sapiens*, existing for around a third of a million years.

A first very basic question within the field of biology, is to ask which species are closest
related to one another. If you think that's an easy question, try to apply it, for example,
on the 8 crocodilian pictures in figure 1.2. You may come up with one or more pairs of
hypothesized closest-related species, but you will probably get stuck on which ancestors
of these pairs are most closely related. Note that doing the other way around, to use
morphology to classify species, is tricky: how to decide when the morphology between
two specimens is different enough to classify these as two species?

Figure 1.2 | Eight members of the order of crocodilians

The second very basic biological question, is to ask *when* these speciation events took place. This question cannot be answered based on morphologies of the present-day species alone, because morphology is a complex trait, and the pace at which morphology changes in time is unknown or unpredictable.

This second question can be answered by using a classical approach, by using the morphology of fossils. This approach can only be used if the species *can* fossilize, and those fossils are found in multiple points in time. Even if this is the case, there are caveats. Using morphology on extinct species is even trickier, as species change their appearance in time. Also an imaginary time machine would not help us out: we could try to determine the number of species in each timepoint, but that would only work if we could confidently define what a species is. We cannot, because speciation is usually a gradual process.

This second question can also be answered using a modern approach, by using the DNA sequences of extant species, as shown, for example, in figure 1.3. Because DNA is inherited from parent to offspring and changes through times, it carries each species' evolutionary histories within it. The point in time when a species speciates is marked by the two daughter species having seperate mutations from that moment on. Due to this, we can easily find closest related species by measuring the similarity in DNA sequences. If we know how frequent mutations occur, we can already do a rough estimation of when the speciation event took place. In reality, DNA sequences of different species varies in length, due to insertions and deletions in genetic sequences, but in the simulation studies in this thesis, we will ignore this.

Figure 1.3 | A 40-nucleotide DNA alignment of six hypothetical species. The species are named A to and including F. The four colors denote the four different nucleotides, in which the red color resembles adenine, yellow depicts cytosine, green is for guanine, and blue resembles thymine. The top row shows the (artificial) root sequence, which is usually unknown.

Using the DNA sequences of extant species to answer our biological questions, however, is a complex topic. For starters, there are multiple ways to do so, like bootstrapping, jack-knifing, parsimony, maximum likelihood and more. A conveniently simple approach is to use UPGMA, which answer our biological questions, in the form of a phylogeny.


Figure 1.4 | Phylogeny created from the alignment in figure 1.3 using UPGMA. This method is irrelevant in the context of this thesis.

Figure 1.4 shows a phylogeny (also called 'phylogenetic tree', or simply 'tree'), created from the alignment in figure 1.3 using one such methodology. It shows the six hypothetical species and their evolutionary relationships. Going from left to right, we travel through time from the past to the present. The leftmost vertical line indicates the first speciation event, which gave rise to the first two ancestral species. This first split in the tree is called the crown, the moment in time this occurred is called the crown age. Sure, we could equally well have started the phylogeny with one ancestral species at the utmost left, going further back in time from the crown age to, what is called, the stem age, but in the context of this thesis, we do not.

The problem with phylogenies is, that it is impossible to go out in the field and measure one, as they depict which species lived when *in the past*. Instead, we *construct* phylogenies. For example, the phylogeny in figure 1.4, how well does match the true phylogeny? **That question, is the main question of this thesis: how well can we construct a phylogeny from an alignment?** What is the error we make when we construct a 'tree without birds'?
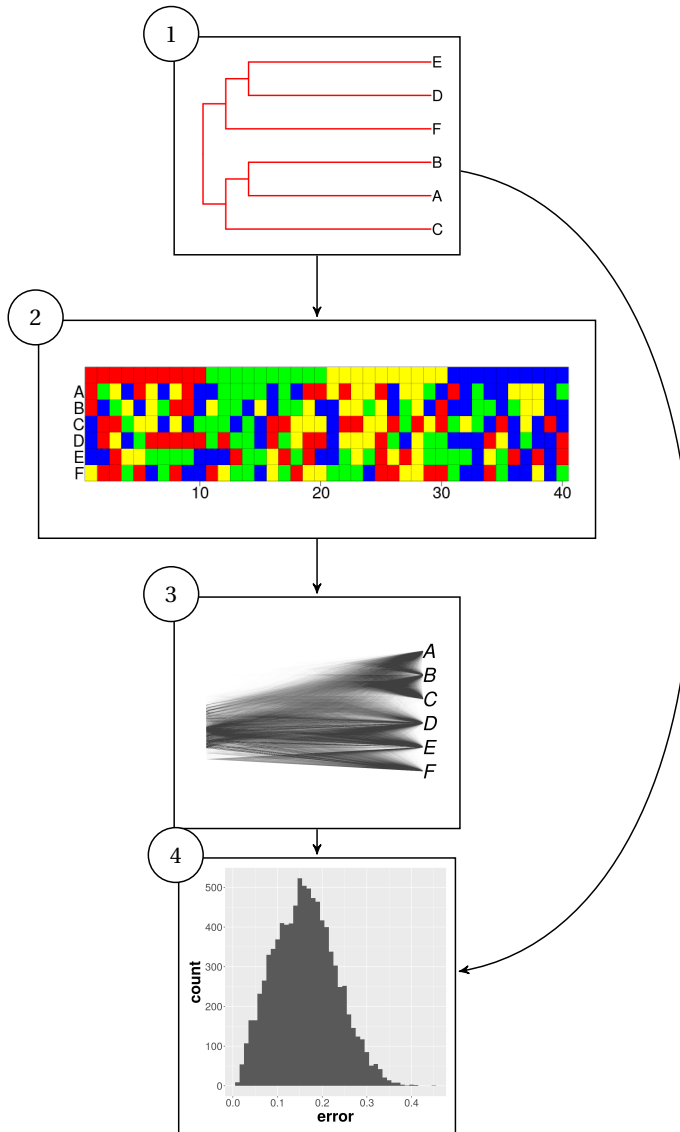
Answering this research question is, at first glance, easy: 1. simulate a true phylogeny. 2. simulate an alignment that follows that phylogeny. 3. construct a phylogeny from that alignment. 4. measure the difference between the true and constructed phylogeny. This workflow is depicted in figure 1.5. All steps, however, are more complex than just this.
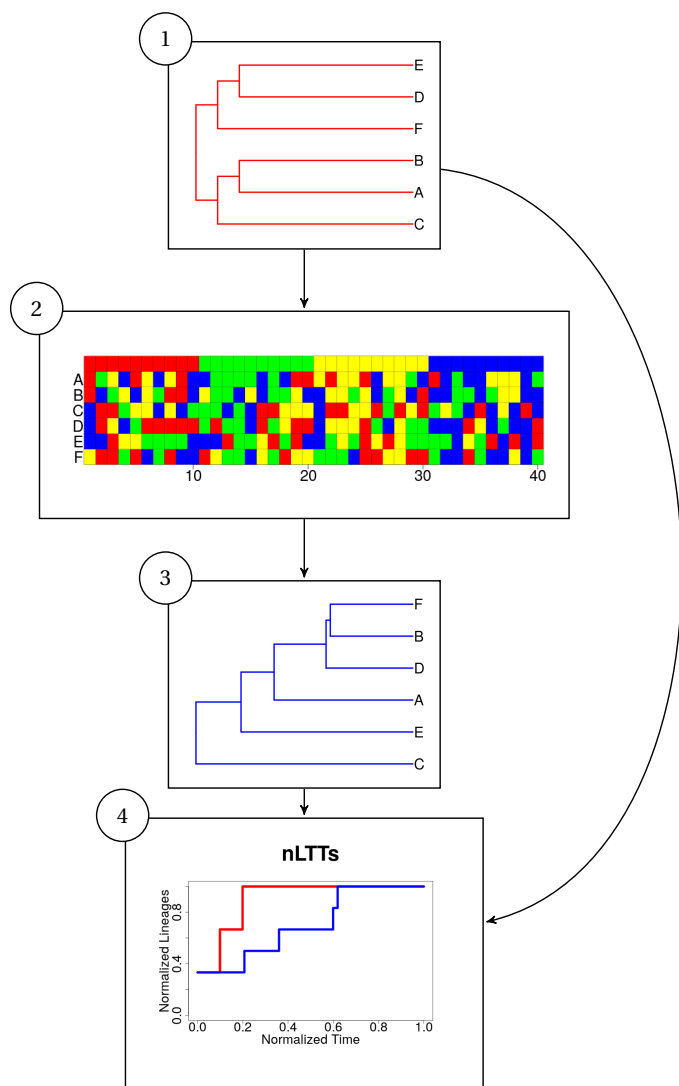
Figure 1.5 | (Simplified) method to answer the research question of this thesis: 1. simulate a true phylogeny. 2. simulate an alignment that follows that phylogeny. 3. construct a phylogeny from that alignment. 4. compare the true and constructed phylogeny.

Constructing a phylogeny from an alignment is the step that gets most attention in this thesis, as it is also the most complex one. Unlike the methods described earlier,

we do not construct one single phylogeny, but we construct a distribution of multiple phylogenies. Within this distribution of multiple phylogenies, the phylogenies that are more likely, will be present more often. This method is called Bayesian phylogenetics, in which we use a Bayesian approach to create phylogenies based on genetics.

Figure 1.6 | Method to answer the research question of this thesis: 1. simulate a true phylogeny. 2. simulate an alignment that follows that phylogeny. 3. infer a distribution of phylogenies from that alignment. 4. compare the true phylogenies with the inferred phylogenies.

With Bayesian inference, we need an alignment and our model assumptions to infer

phylogenies. The model assumptions specify what we assume to be true regarding how the alignment came to be. For example, we can assume that 1. the true phylogeny had a constant speciation and extinction rate, 2. the mutation rate is constant and equal for all species and 3. all mutations between nucleotides are equally likely.

The result of a Bayesian inference, is -to be precise- a posterior disribution of jointly-inferred phylogenies and model parameter estimates, simply called 'posterior' in this thesis. The way such a posterior is generated assures that more likely phylogenies are present more often. This distribution of phylogenies shows the (un)certainty of the inference. For example, the posterior phylogenies in figure 1.6, panel 3, show a high degree of uncertainty, as the inferred phylogenies vary widely in shape. The posterior correctly suggests two clades (ABC and DEF), but does not confidently show the two most related taxa (AB and DE). We can already make the rough claim that, would the phylogeny in panel 1 in figure 1.6 depict a true phylogeny, we make a big error in its inference.

To be able to do the phylogentic inference needed for the rest of this thesis, I developed an R package to do so, which is discussed in chapter 2. Because a Bayesian inference can be set up in many ways, the greatest asset of that package is that it gives a consistent grammar to express each setup. Additionally, the R package allows to run Bayesian inference from the command-line, which is essential for the theoretical studies in this thesis.

If we can measure the error we make in our inference, we can try and improve the inference. One way to improve it, is to use a better inference model. Ideally, we would use the same inference model that gave rise to the true phylogeny and alignment, but, alas, we (usually) do not know that model. We do not know that model, because we do not know the model that nature used: the processes that cause speciation are possibly many, and the mechanisms of each are unknown and/or debatable. Due to this, we'll have to resort to the many phylogentic models to explain the DNA (RNA, protein, morphological and fossil) data best.

There are plenty of phylogentic models, ranging from simplistic to very complex. The most popular models make it into our phylogenetic programs (which, in turn, may make these models even more popular), which I will define as 'standard models'. When empiricist build a phylogenetic tree from their painstakingly acquired DNA alignment, they pick their favorite standard model or use an algorithm to select one. The empiricist assumes that the standard models are good enough for his/her cause.

The standard models, however, make some assumptions that will not hold in all biological cases. This will increase the error we make in our inference. But will that error be profound enough to reject using a standard model?

To be able to determine the impact of using a standard phylogentic model, when we know the biological process in more complex than it assumes, me and Giovanni Laudanno developed an R package to quantify the error we make due to this mismatch, which is described in chapter 3.

One assumption of all standard models is that speciation event happen independently, that is, there are never two speciation events at the same time. There are biological scenario's in which may say this is false: when a habitat is split up, due to a geological barrier, this will result in two species communities. The change from one to two communities is likely to affect both communities and trigger a speciation event in both communities.

The inference error of ignoring co-occurring speciation is quantified by me and Giovanni Laudanno in chapter 4.

Another assumption of all standard models is that speciation event happen instantaneously, that is, when there is a speciation event, the two species are immediatly recognized as such. We know that speciation takes time. The inference error of ignoring this fact is quantified by me in chapter 5.

In chapter 6, I show which conclusions can be drawn from these chapters

1.0.1. PHOTO ATTRIBUTION

Figure 1.1, the left image, Preserved specimen of chalumnae by Alberto Fernandez Fernandez is licensed under CC BY-SA 3.0. the image at the right, Akha couple in northern Thailand by Weltenbummler84 is licensed under CC BY-SA 2.0 DE.

For figure 1.2, the selection of these eight images was done by OneZoom. These images, from top-left to bottom-row, row-first: Nile Crocodile by Marco Schmidt is licensed under CC-BY-SA 3.0. American Alligator by NASA Kennedy. False Gavial by Yinan Chen is marked as being in the public domain. Common Caiman by Michael Wolf is licensed under CC-BY 2.5. African Dwarf Crocodile by Staycoolandbegood is marked as being in the public domain. Smooth-fronted Caiman by Whaldener Endo is licensed under CC-BY-SA 4.0 Phillipines Crocodile by Vanderploeg is marked as being in the public domain. Board-snouted Crocodile by Hadonos is Marked as being in the public domain.